



UNIVERSITY OF
CAMBRIDGE

OCaml Workshop: Oxford, UK, September 8th, 2017.

Owl - A General-Purpose Numerical Library in OCaml

Liang Wang

University of Cambridge


September 2017

Motivation

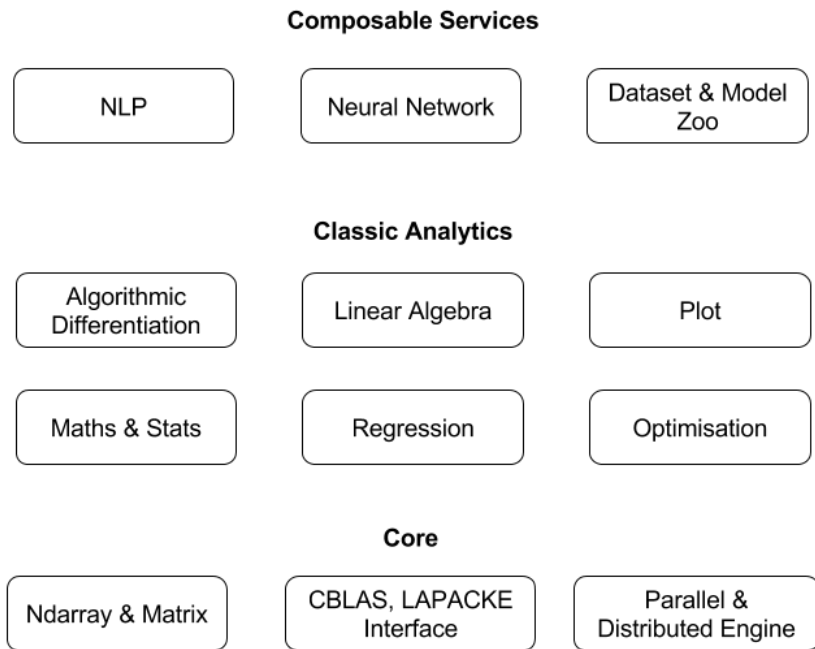
- How can we attract more people into OCaml community?
- OCaml core library is kept minimal (on purpose) for ?
- But many applications need strong numerical supports.
- Libraries for numerical computing in OCaml are fragmented.
- Difference between system library and numerical library:
 - A **reduced** set of operations atop of system abstraction.
 - A **complex** set of functions atop of number abstraction.

Highlights

- Owl is an experimental project to develop a numerical platform.
- One year intensive development: **> 130k LOC, 6.5k functions**.
- A comprehensive set of classic numerical functions.
- Strong support for modern data analytics (ML & DNN).
- As fast as C but as concise as Python with static type-checking.



Architecture



- **Composable Services** layer enables fast development of modern data analytics using type-safe functional programming.
- **Classic Analytics** layer implements a comprehensive set of classic analytic functions such as various hypothesis tests.
- **Core** layer provides the building block of the whole numerical system and the capability of distributed computing.

Ndarray - Module Structure

- `Dense` and `Sparse` as two top modules.
- `S/D/C/Z` submodules to deal with four number types

```
Dense.Ndarray.S.zeros [|5;5|];;      (* single precision real ndarray *)
Dense.Ndarray.D.zeros [|5;5|];;      (* double precision real ndarray *)
Dense.Ndarray.C.zeros [|5;5|];;      (* single precision complex ndarray *)
Dense.Ndarray.Z.zeros [|5;5|];;      (* double precision complex ndarray *)
```

- There is also an `Any` module to deal with all other types.

```
let x = Dense.Ndarray.Any.create [|5;5;5|] true in
  Dense.Ndarray.Any.get_slice_simple [|-1;0|]; [|]; [|] x;;
```



Ndarray - Groups of Functions

- Creation functions: `zeros`, `gaussian`, `linspace`, `bernoulli`, `magic` ...
- Manipulation functions: `pad`, `slice`, `flip`, `rotate`, `transpose`, `tile` ...
- Iteration functions: `iter`, `map`, `filter`, `fold`, `map2`, `exists`, `for_all` ...
- Comparison functions: `equal`, `greater`, `less`, `elt_*`, `*_scalar` ...
- Vectorised uni- and binary maths functions: `sin`, `tanh`, `log`, `fix`, `elu` ...

Ndarray - Performance Critical Code

```
#define FUN1 real_float_is_negative
#define NUMBER float
#define STOPFN(X) (X >= 0)
#include "owl_dense_common_vec_cmp.c"

#define FUN1 real_double_is_negative
#define NUMBER double
#define STOPFN(X) (X >= 0)
#include "owl_dense_common_vec_cmp.c"

#define FUN1 complex_float_is_negative
#define NUMBER complex_float
#define STOPFN(X) (X.r >= 0 || X.i >= 0)
#include "owl_dense_common_vec_cmp.c"

#define FUN1 complex_double_is_negative
#define NUMBER complex_double
#define STOPFN(X) (X.r >= 0 || X.i >= 0)
#include "owl_dense_common_vec_cmp.c"
```

```
#define FUN4 real_float_abs
#define NUMBER float
#define NUMBER1 float
#define MAPFN(X) (fabsf(X))
#include "owl_dense_common_vec_map.c"

....

#define FUN5 real_float_sum
#define INIT float r = 0.
#define NUMBER float
#define ACCFN(A,X) (A += X)
#define COPYNUM(X) (caml_copy_double(X))
#include "owl_dense_common_vec_fold.c"
```



Ndarray - Indexing & Slicing

The most fundamental operation; the key to concise code.

- In Numpy : `x[1, 0:5; -1:0:-2]`
- In Julia : `x[1, 0:5, 1:end]`
- In Owl : `get_slice_simple [[1]; [0;5]; [-1;0;-2]] x`

Ndarray - More Complex Slicing

```
type index =  
  | I of int          (* single index *)  
  | L of int list     (* list of indices *)  
  | R of int list     (* index range *)  
  
type slice = index list  
  
val get_slice : index list -> ('a, 'b) t -> ('a, 'b) t  
val set_slice : index list -> ('a, 'b) t -> ('a, 'b) t -> unit
```

E.g., [I 5; L [2;1;0]; R [2;-1;2]]

Ndarray - Operator Design

Commonly used operators have been implemented: `+` `-` `*` `/` `+$`
`-$` `=` `.=` `=~` `>` `>.` `>$` `...`

Operators are divided into groups.
E.g., Ndarray does not support `*@`

Operators are implemented as functors to avoid maintaining multiple copies of definition.

We can easily change the definition for the whole system easily in one place.

```
module Make_Basic (M : BasicSig) = struct
  type ('a, 'b) op_t0 = ('a, 'b) M.t
  let ( + ) = M.add
  let ( - ) = M.sub
  let ( * ) = M.mul
  let ( / ) = M.div
  let ( +$ ) = M.add_scalar
  ...
end
```

```
module Operator = struct
  include Owl_operator.Make_Basic (Owl_dense_matrix_generic)
  include Owl_operator.Make_Extend (Owl_dense_matrix_generic)
  include Owl_operator.Make_Matrix (Owl_dense_matrix_generic)
end
```

Ndarray - Challenges

- Interoperation on different numbers - `Ext` module.

```
Ext.(F 10. * Dense.Matrix.C.uniform 5 5 >. F 5.)
```

- How about a generic abs using current Bigarray types?

```
external owl_real_float_min_i : int -> ('a, 'b) owl_vec -> int = "real_float_min_i"
external owl_real_double_min_i : int -> ('a, 'b) owl_vec -> int = "real_double_min_i"
external owl_complex_float_min_i : int -> ('a, 'b) owl_vec -> int = "complex_float_min_i"
external owl_complex_double_min_i : int -> ('a, 'b) owl_vec -> int =
"complex_double_min_i"
```

```
let _owl_min_i : type a b. (a, b) kind -> (a, b) owl_vec_op01 = function
| Float32    -> owl_real_float_min_i
| Float64    -> owl_real_double_min_i
| Complex32  -> owl_complex_float_min_i
| Complex64  -> owl_complex_double_min_i
| _          -> failwith "_owl_min_i: unsupported operation"
```



CBLAS/LAPACK Interface

- Not **BLAS/LAPACK** because Owl sticks to **C-layout**.
- Interface is automatically generated from C header files.
- Rely on **Ctypes**, but only used its c stub file (**foreign** is an issue).
- The challenge is the versioning and deployment on different platforms.

CBLAS/LAPACKE Interface to Linalg

- Automatically generated code only provides a thin wrapping.
- Still need hand-writing interface to wrap up $S/D/C/Z$ types.
- How much to expose to Linalg? Flexibility vs. Convenience.
- It is often believed that C-interface introduces some overhead,
but Owl often achieves better performance than Julia.

Distributed Computing - Design Rationale

Think about **distributed analytics** -> it is **distribution + analytics**

- Parallel & distributed engine must be separated out.
- Engine APIs must be minimal and simple for easy composition.
- Able to deal with both low-level and high-level data structures.
- Avoid letting developers deal with details like message passing.

Owl + Actor (Sub)System

- Three engines: [Map-Reduce](#); [Parameter Server](#); [Peer-to-Peer](#)
- Specifically designed synchronisation mechanism for scalability.

```
val start : ?barrier:barrier -> string -> string -> unit
val register_barrier : ps_barrier_typ -> unit
val register_schedule : ('a, 'b, 'c) ps_schedule_typ -> unit
val register_pull : ('a, 'b, 'c) ps_pull_typ -> unit
val register_push : ('a, 'b, 'c) ps_push_typ -> unit
val register_stop : ps_stop_typ -> unit
```

```
type barrier =
  | ASP      (* Asynchronous Parallel *)
  | BSP      (* Bulk Synchronous Parallel *)
  | SSP      (* Stale Synchronous Parallel *)
  | PSP      (* Probabilistic Synchronous Parallel *)
```

Owl + Actor : Neural Network Example

```
let network =  
  input [|28;28;1|]  
  |> lambda (fun x -> Maths.(x / F 256.))  
  |> conv2d [|5;5;1;32|] [|1;1|] ~act_typ:Activation.Relu  
  |> max_pool2d [|2;2|] [|2;2|]  
  |> dropout 0.1  
  |> fully_connected 1024 ~act_typ:Activation.Relu  
  |> linear 10 ~act_typ:Activation.Softmax  
  |> get_network
```

A convolutional neural network can be defined as concise as that in the state-of-the-art system specialised in deep neural networks.

Changing it into a distributed algorithm just requires one line of code thanks to Owl's cutting edge parallel & distributed computation engine.

Make it distributed! Yay!

```
module M2 = Owl_neural_parallel.Make ( Owl.Neural.S.Graph )  
                                     (Actor.Param)
```


Owl + Actor : Ndarray Example

Similarly, we can also transform a Ndarray into a distributed Ndarray, by simply

```
module M2 = Owl_parallel.Make ( Dense.Ndarray.S ) (Actor.Mapre)
```

Composed by a functor in Owl_parallel module, which connects two systems and hides details.

```
module type Mapre_Engine = sig
  val map : ('a -> 'b) -> string -> string
  val map_partition: ('a list -> 'b list) -> string -> string
  val union : string -> string -> string
  val reduce : ('a -> 'a -> 'a) -> string -> 'a option
  val collect : string -> 'a list
  val workers : unit -> string list
  val myself : unit -> string
  val load : string -> string
  val save : string -> string -> int
end
```

```
module type Ndarray = sig
  val shape : arr -> int array
  val empty : int array -> arr
  val create : int array -> elt -> arr
  val zeros : int array -> arr
  ...
end
```



Algorithmic Differentiation

- Core component to bridge the gap between low-level numerical functions and high-level analytical models.
- Functor to support both single- and double-precision.
- Support quite many functions and operations already.

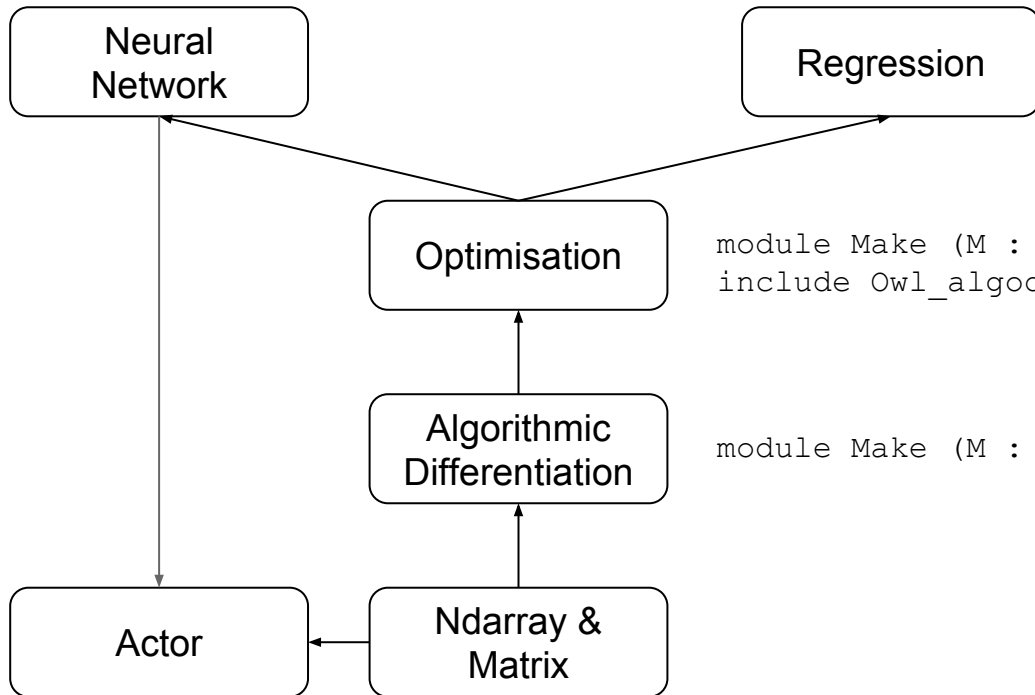
```
(* AD module of Float32 type *)  
module S = Owl_algodiff_generic.Make (Owl_dense_matrix.S) (Owl_dense_ndarray.S)  
  
(* AD module of Float64 type *)  
module D = Owl_algodiff_generic.Make (Owl_dense_matrix.D) (Owl_dense_ndarray.D)
```



Revisit Owl's Architecture

```
module Make (M : ...) (A : ...) = struct
include Owl_optimise_generic.Make (M) (A)
```

```
module Make (M : ...) (A : ...) = struct
include Owl_optimise_generic.Make (M) (A)
```



```
module Make (M : MatrixSig) (A : NdataArraySig) = struct
include Owl_algodiff_generic.Make (M) (A)
```

```
module Make (M : MatrixSig) (A : NdataArraySig)
```

Regression - OLS

```
let ols ?(i=false) x y =  
  let params = Params.config  
    ~batch:(Batch.Full) ~learning_rate:(Learning_Rate.Adagrad 1.)  
    ~gradient:(Gradient.GD) ~loss:(Loss.Quadratic) ~verbosity:false  
    ~stopping:(Stopping.Const 1e-16) 1000.  
  in  
  _linear_reg i params x y
```

Regression - Ridge

```
let ridge ?(i=false) ?(a=0.001) x y =  
  let params = Params.config  
    ~batch:(Batch.Full) ~learning_rate:(Learning_Rate.Adagrad 1.)  
    ~gradient:(Gradient.GD) ~loss:(Loss.Quadratic)  
    ~regularisation:(Regularisation.L2norm a) ~verbosity:false  
    ~stopping:(Stopping.Const 1e-16) 1000.  
  in  
  _linear_reg i params x y
```



Regression - SVM

```
let svm ?(i=false) ?(a=0.001) x y =  
  let params = Params.config  
    ~batch:(Batch.Full) ~learning_rate:(Learning_Rate.Adagrad 1.)  
    ~gradient:(Gradient.GD) ~loss:(Loss.Hinge)  
    ~regularisation:(Regularisation.L2norm a) ~verbosity:true  
    ~stopping:(Stopping.Const 1e-16) 1000.  
  in  
  _linear_reg i params x y
```



Neural Network

- Built atop of `Algodiff` Module, `pros > cons`
- Support both simple Feedforward and Graph structure.
- Many neurons have been implemented: `linear`, `conv2d`, `maxpool ...`
- Code is as concise as state-of-the-art specialised library.
- Performance can be improved by Actor system.

Neural Network - VGG Example

```
let make_network input_shape =  
  input input_shape  
  |> normalisation ~decay:0.9  
  |> conv2d [|3;3;3;32|] [|1;1|] ~act_typ:Activation.Relu  
  |> conv2d [|3;3;32;32|] [|1;1|] ~act_typ:Activation.Relu ~padding:VALID  
  |> max_pool2d [|2;2|] [|2;2|] ~padding:VALID  
  |> dropout 0.1  
  |> conv2d [|3;3;32;64|] [|1;1|] ~act_typ:Activation.Relu  
  |> conv2d [|3;3;64;64|] [|1;1|] ~act_typ:Activation.Relu ~padding:VALID  
  |> max_pool2d [|2;2|] [|2;2|] ~padding:VALID  
  |> dropout 0.1  
  |> fully_connected 512 ~act_typ:Activation.Relu  
  |> linear 10 ~act_typ:Activation.Softmax  
  |> get_network
```

Define a network structure in a very concise and functional way.

Neural Network - Cost of ...

```
module Normalisation = struct

  let run x l =
    let a = F (1. /. float_of_int (shape x).(l.axis)) in
    l.mu <- Maths.(a * (sum_ ~axis:l.axis x));
    l.var <- Maths.(a * (sum_ ~axis:l.axis (x * x)));
    let x' = Maths.((x - l.mu) / sqrt (l.var + F 1e-8)) in
    Maths.(x' * l.gamma + l.beta)

  ...

end
```

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

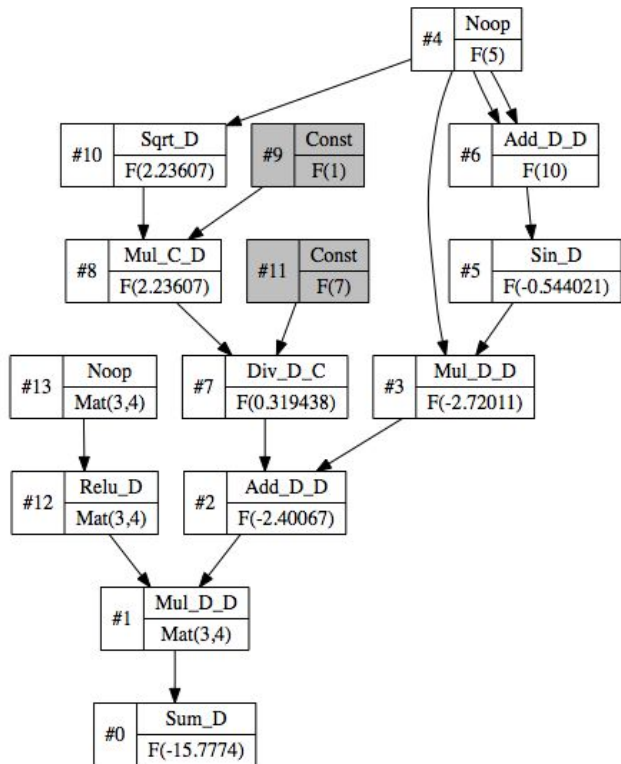
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Adding a new neuron type is simple, [Algodiff](#) will take care of calculating derivatives in the backpropagation phase.

Computation Graph - Simple Function

```
let f x y = Maths.((x * sin (x + x) + ( F 1. * sqrt x) / F 7.) * (relu y) |> sum)
```




Debugging can be tricky, and visualisation is often very helpful.

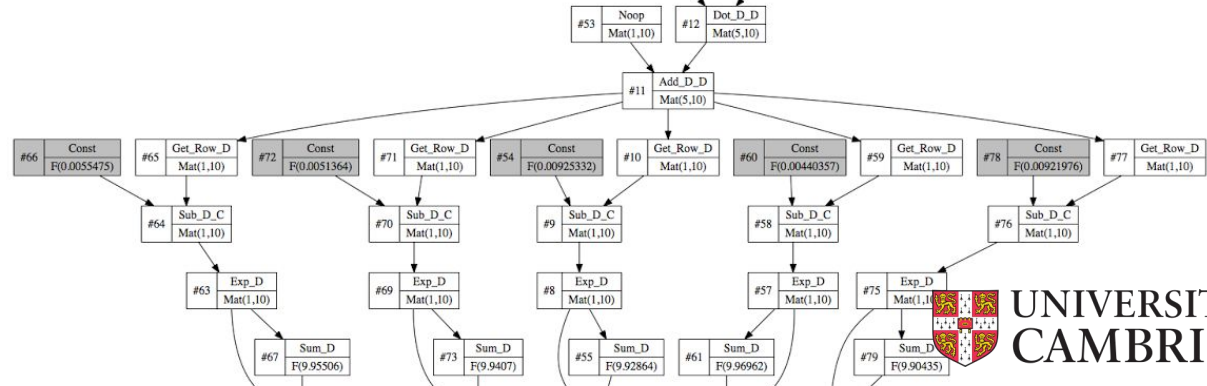
Owl can print out raw trace on terminal in human-readable format, or print in dot file format for further visualisation.

The node contains rich information for debugging.

Computation Graph - VGG Network

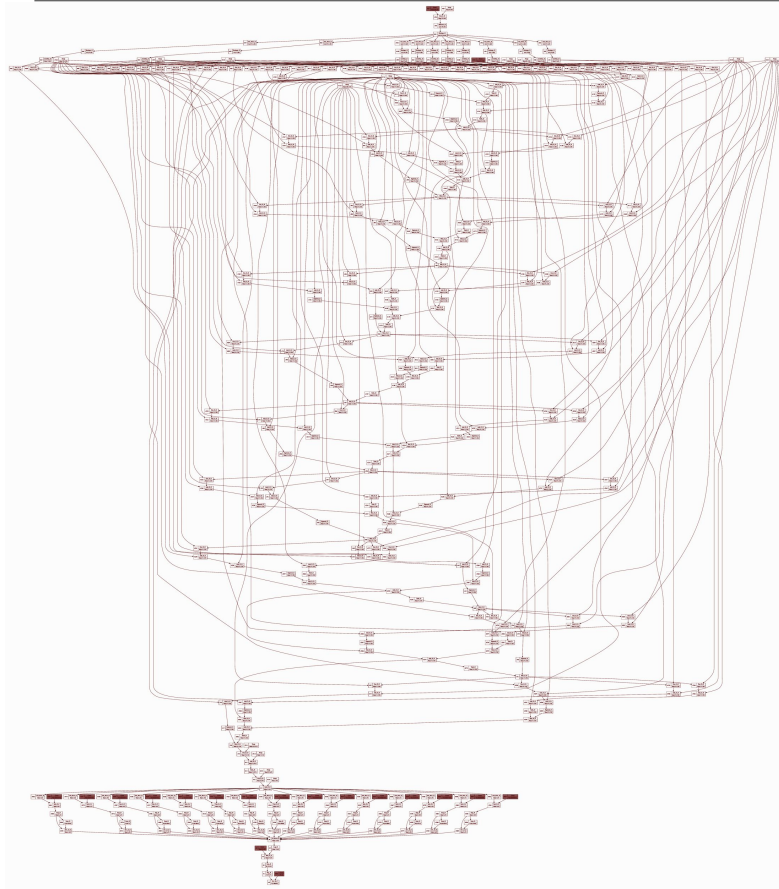


```
input input_shape
|> normalisation
|> conv2d [|3;3;3;32|] [|1;1|] ~act_typ:Activation.ReLu
|> conv2d [|3;3;32;32|] [|1;1|] ~act_typ:Activation.ReLu
~padding:VALID
|> max_pool2d [|2;2|] [|2;2|] ~padding:VALID
|> dropout 0.1
|> conv2d [|3;3;32;64|] [|1;1|] ~act_typ:Activation.ReLu
|> conv2d [|3;3;64;64|] [|1;1|] ~act_typ:Activation.ReLu
~padding:VALID
|> max_pool2d [|2;2|] [|2;2|] ~padding:VALID
|> dropout 0.1
|> fully_connected 512 ~act_typ:Activation.ReLu
```



UNIVERSITY OF
CAMBRIDGE


Computation Graph - LSTM Network



```
input [|wndsz|]  
|> embedding vocabsz 40  
|> lstm 128  
|> linear 512 ~act_typ:Activation.ReLu  
|> linear vocabsz ~act_typ:Activation.Softmax  
|> get_network
```



Google Inception-V3




UNIVERSITY OF
CAMBRIDGE

Plotting Functions

- Built atop of `ocaml-plplot`, but with improved interface.
- Hide all the dirty details in `Plplot`, provide matlab-like APIs.
- The core plotting engine is very small **< 200 LOC**

```
let h = Plot.create ~m:2 ~n:3 "plot_017.png" in
Plot.subplot h 0 0;
Plot.surf ~h x y z;
Plot.subplot h 0 1;
Plot.mesh ~h x y z;
Plot.subplot h 0 2;
Plot.(surf ~h ~spec:[ Contour ] x y z);
Plot.subplot h 1 0;
Plot.(mesh ~h ~spec:[ Contour; Azimuth 115.; NoMagColor ]
Plot.subplot h 1 1;
Plot.(mesh ~h ~spec:[ Azimuth 115.; ZLine X; NoMagColor;
Plot.subplot h 1 2;
Plot.(mesh ~h ~spec:[ Azimuth 115.; ZLine Y; NoMagColor; RGB (150,255,40) ] x y z);
Plot.output h;;
```



Zoo System - Share Code Snippets

- Originally designed for sharing neural network models.
- Now used for sharing small code snippet and scripting in Owl.
- Enhance Owl's numerical capability and fast prototyping.
- Very simple design: add one more directive `#zoo` to toplevel.

Zoo System - Example

gist.github.com/ryanrhymes/2e7c902812a7ae0547e24f7ea743c7e6

当前 BigData DSP Network Theory Unsupervised game theory machine learning stochastic reading hadoop liang misc

GitHubGist Search... All gists GitHub New gist

ryanrhymes / cifar10.ml Secret

Last active a day ago

Code Revisions 5 Embed <script src="https://gi... Download ZIP

CIFAR10 CNN Example

readme.md Raw

CIFAR10 VGG Example

This example demonstrates how to build a VGG-like convolutional neural network for CIFAR10 dataset.

cifar10.ml Raw

```
1 #!/usr/bin/env owl
2
3 open Owl
4 open Owl_neural
5 open Algodiff.S
6 open Owl_neural_neuron
7
8
9 let model () =
10   let open Owl_neural_graph in
11   let nn = input [32;32;3]
12   |> conv2d [3;3;3;32] [1;1] ~act_type:Activation.Relu
13   |> conv2d [3;3;32;32] [1;1] ~act_type:Activation.Relu ~padding:Owl_dense_ndarray_generic.VALID
14   |> max_pool2d [2;2] [2;2] ~padding:Owl_dense_ndarray_generic.VALID
15   |> dropout 0.1
16   |> conv2d [3;3;32;64] [1;1] ~act_type:Activation.Relu
17   |> conv2d [3;3;64;64] [1;1] ~act_type:Activation.Relu ~padding:Owl_dense_ndarray_generic.VALID
```

Publish on gist; then use in your script with `#zoo` directive.

```
#!/usr/bin/env owl

#zoo "2e7c902812a7ae0547e24f7ea743c7e6"
#zoo "217ef87bc36845c4e78e398d52bc4c5b"
```

Owl's Zoo System

Usage:

```
owl [utop options] [script-file]
owl -upload [gist-directory]
owl -download [gist-id]
owl -remove [gist-id]
owl -update [gist-ids]
owl -run [gist-id]
owl -info [gist-ids]
owl -list
owl -help
```



UNIVERSITY OF
CAMBRIDGE

Future Plan

- There is a great space for further optimisation.
- C code can be further optimised with SSE ...
- Actor will include an engine for multicore OCaml.
- Matrix and Nddarray can be unified, leads to simpler code.
- Indexing needs enhancement, ppx, performance, and etc.

Thank you!

Website: <https://github.com/ryanrhymes/owl>

....

- How to architect a modern numerical library
- Meet modern need.
- Extensible, Flexible, Manageable
- Focus on architectural design, avoid doing thing that can be done by the machine: generate interface, calculating derivative
- Do as much as we can with as little code as possible.
- Why ... reduce potential error.